



SE350 RTX LAB 2

P2-A Interprocess Communications (IPC)

P2-B Timing Service

Irene Yiqing Huang

Electrical and Computer Engineering Department

P2-B Timing Service

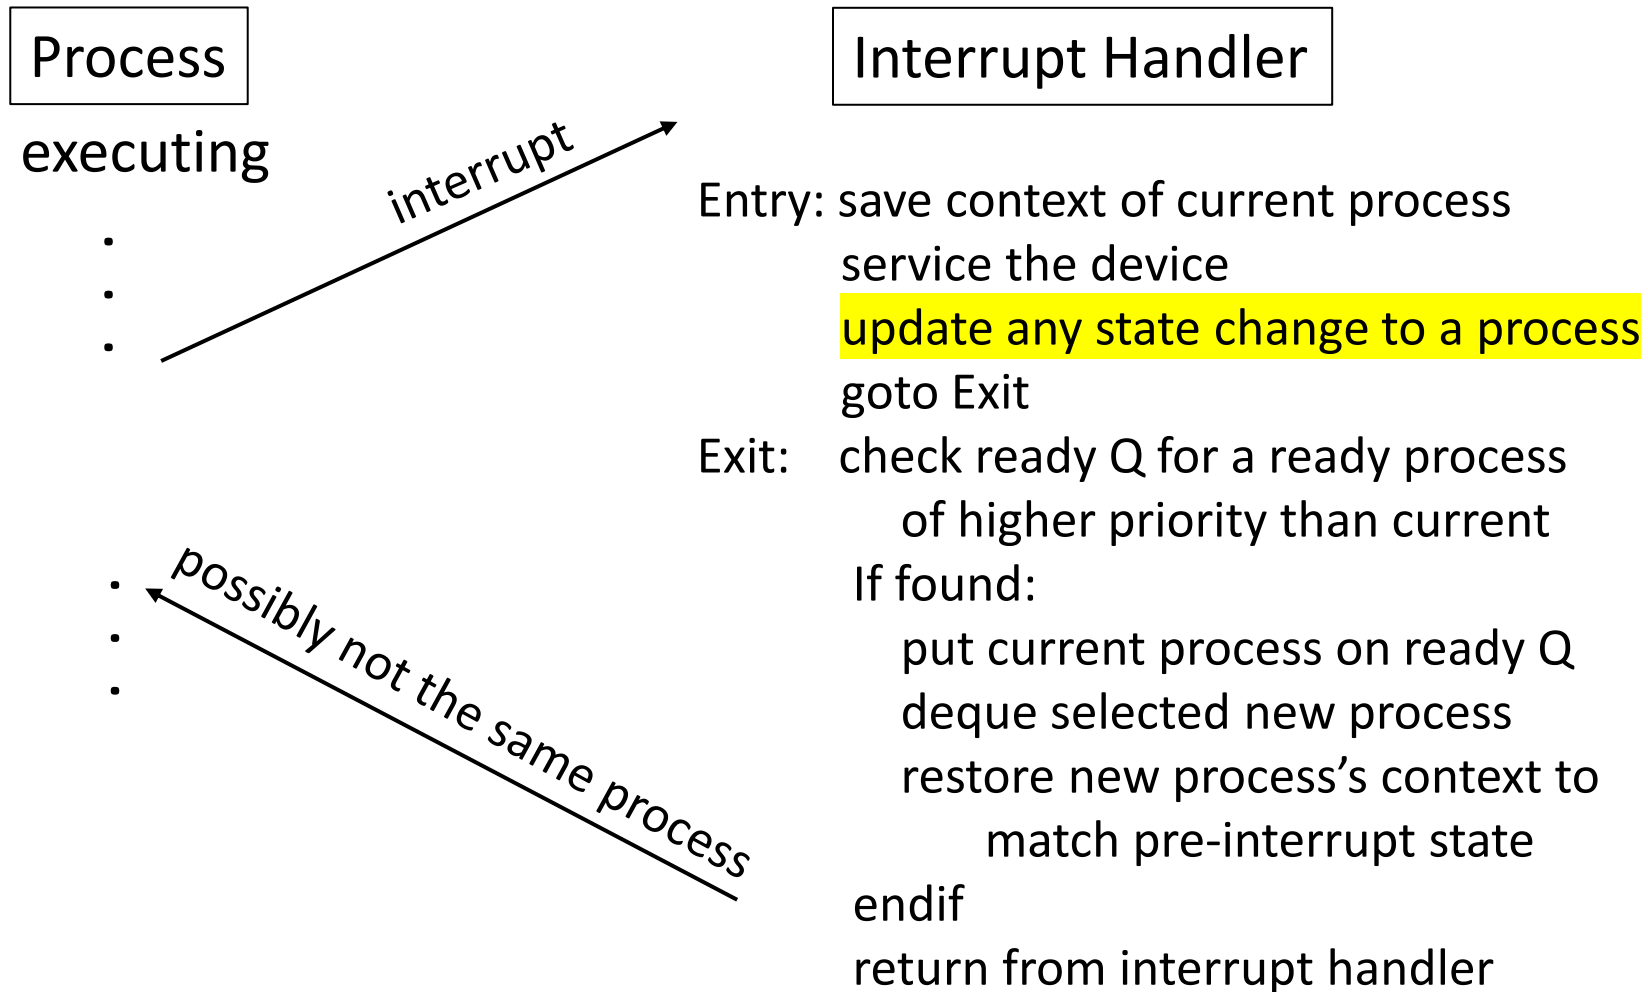
Interrupt Handling

- In real-time OS, interrupt handling must be fast
 - short latency to respond to interrupt
 - fast processing by interrupt handler
- Interrupts may cause a change in state for some blocked process
 - e.g. a process blocked for external event to occur will have its state changed to ready and be placed on the ready process queue when the event occurs

Interrupt Handling Issues

- Possible interpretation of an interrupt:
 - an interrupt is a *hardware message* usually requiring a short latency and quick service
- Design issues:
 - does the interrupt handling code run as part of kernel, within a process (if yes, which?)
 - are interrupt handlers themselves interruptible?
 - if OS is preemptive, need to deal with the possibility that an interrupt results in a higher priority process becoming ready

Interrupt Handling Sequence




Interrupt Handler Design

- Interrupt handler must interact with OS processes
- An **i-process** gets the CPU from an interrupt handling sequence, not through the scheduler.
- It **never blocks** if it invokes a kernel primitive
- The interrupt (exception) handling routine starts the appropriate i-process.
- Conceptually: i-process has the highest priority, is **scheduled by interrupt**.
- No nested interrupt handling

I-process

- State diagram for an i-process:



- A PCB is associated with each i-process with a special state = i-process (permanently)
- It is always ready to run, but not on any ready Q

I-process: Constraints

- An i-process can invoke kernel primitives:
 - however, an i-process is not allowed to block!
- Primitives which can block a process **must** be modified to ensure that **i-process does not block!**
 - e.g. the synchronous receive message primitive
 - return null if the invoking process is an i-process and there is no message waiting
 - similarly for other primitives

I-process: Example

```
__asm void TIMERO_IRQHandler {  
    atomic(on)  
    save the context of the current_process ;  
    switch the current_process with timer_i_process ;  
    load the timer_i_process context ;  
    call the timer_i_process C function ;  
    invoke the scheduler to pick next to run process ;  
    restore the context of the newly picked process ;  
    atomic(off)  
    return  
}
```

- Embedded assembly programming
- Context_Switching_IRQ project in the starter code on github can be modified to become a Timer i-process (add interrupt off/on et. al..).

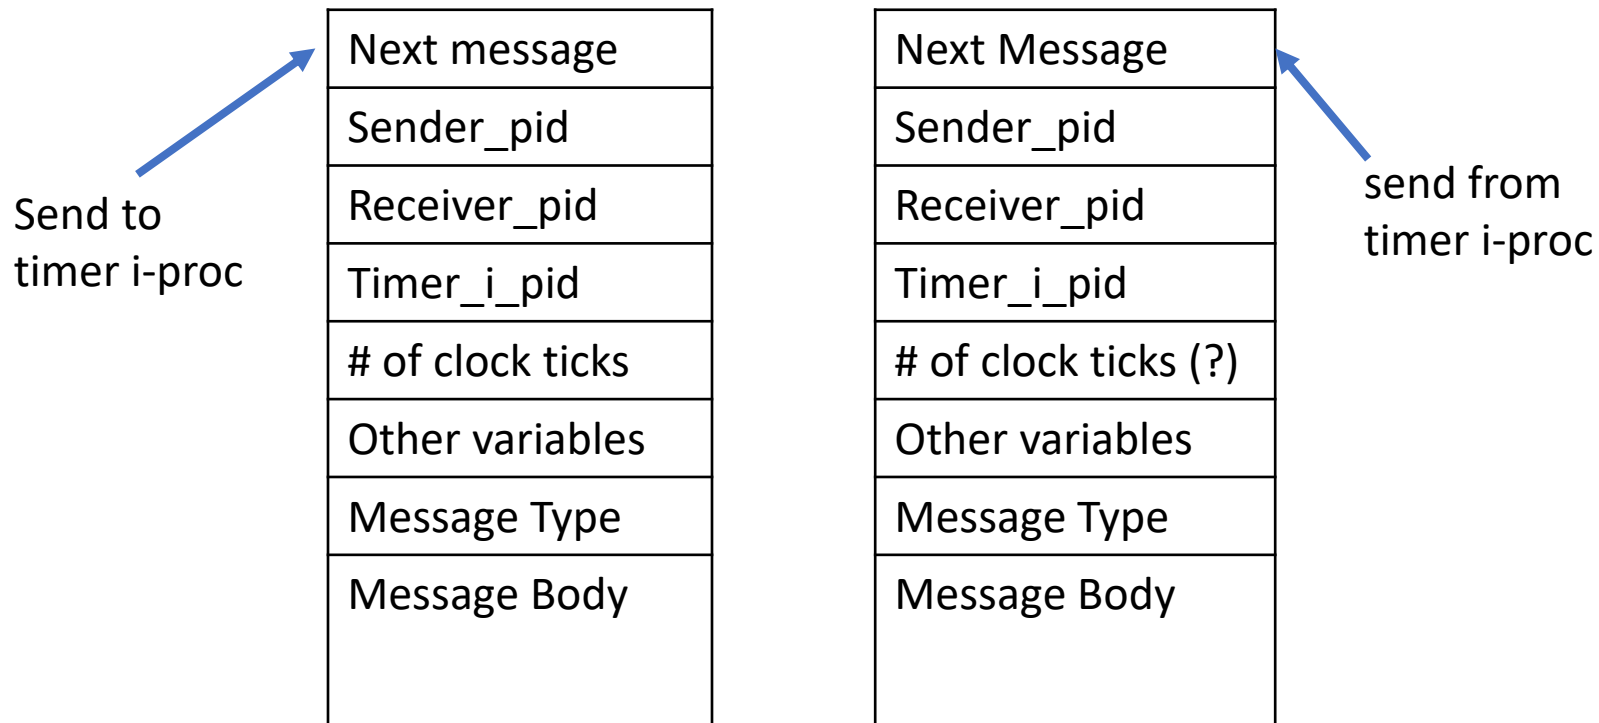
Timing Services: Design

- Fundamental to Real-time OS
- Two parts: interface/protocol design, internal design
- Interface/protocol design
 - basic service only (timeout), no cancellation
 - service request, expiry notification: by messages
- Internal design
 - timing service implemented by *i-process*
 - service request: a user process sends a request message to the timing i-process
 - timeout notification: the i-process sends message back

Timing Services: Request

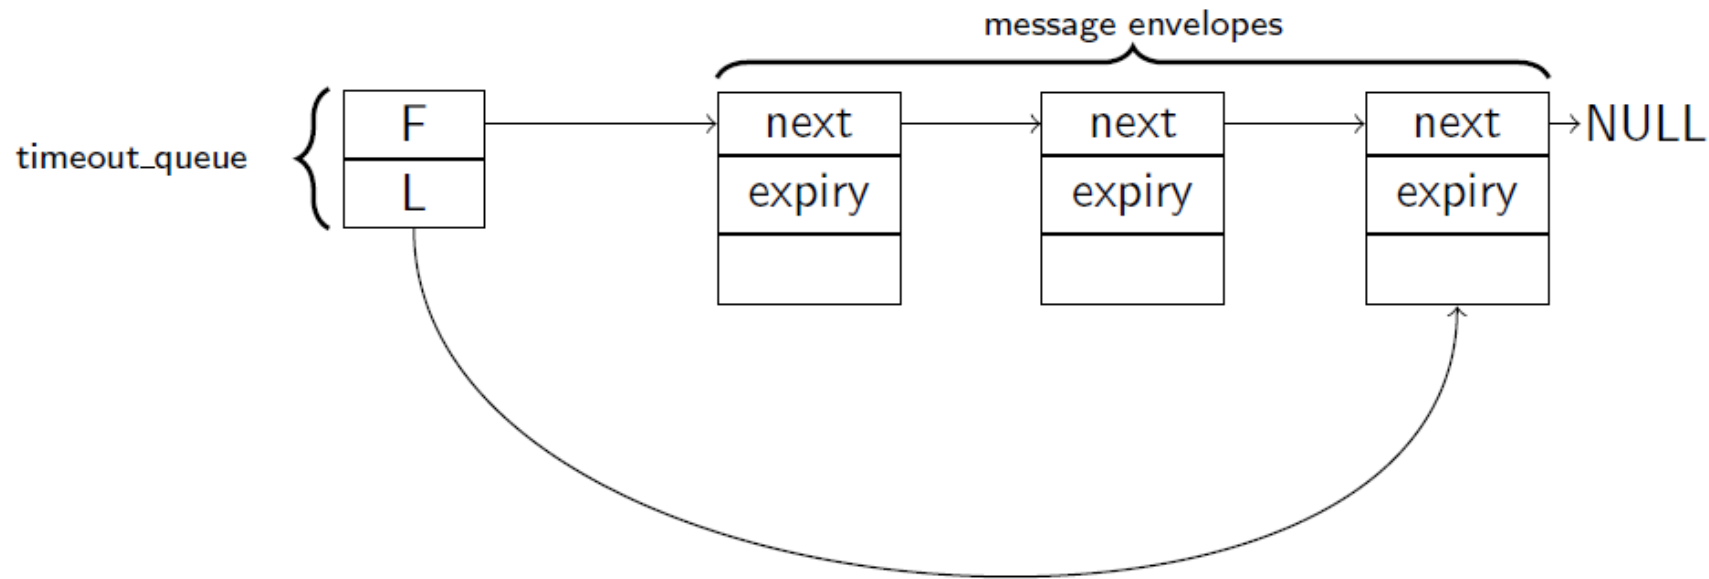
```
int delayed_send(int rcv_pid, void *msg_env, int delay)
```

- The timer i-process receives messages to deliver with a time-out value (i.e. the delay).



Timing Services: Request

- After the delay expires, the timer i-process forwards the message to the receiver
- The timer i-process holds requests in a sorted list.



Timing Services: Expiry Time

- To reduce CPU overhead, expiry time can be replaced by the # of clock ticks after the expiry of the predecessor in the list
- Example: queue timeout list {25, 30, 0, 10}
 - one timeout for 25 clock ticks
 - two timeouts for 55 clock ticks
 - one timeout for 65 clock ticks

Timer I-process

At each clock tick (i.e., interrupt), the timer i-process:

- Increments current time
- Calls receive message repeatedly to retrieve new requests (non-blocking)
- If there are new requests, it adds them to the queue (maintaining sorted order)
- Checks if any timing requests have expired
 - If yes, send the message to the destination

Timer I-process

```
void timer_i_process ( ) {  
    update timeout in queue  
    /* get pending requests */  
    while ( pending messages to i-process ) {  
        insert envelope into the timeout queue ;  
    }  
    while ( first message in queue timeout expired ) {  
        env ← dequeue ( timeout_queue ) ;  
        target_pid ← destination_pid from env;  
        /* forward msg to destination */  
        k_send_message ( target_pid , env ) ;  
    }  
}
```

Design for Preemption

- On return from send() or any primitives that may make another process ready
 - Check if the highest priority ready process's priority is greater than the current process priority
- What the current process is i-process?
- Possible solution
 - Let priority of i-process be the highest
 - On return from interrupt, check whether the priority of the interrupted process still the highest
 - If not, context switch to the higher priority ready process

RTX Initialization

- What operations need to be carried out at start-up?
- Initialize all hardware, incl.
 - Board system Initialization
 - Interrupts (hardware and software: vector table & traps)
 - **Timer(s)** and ~~Serial port(s)~~
- Create all kernel data structures
 - Memory management kernel data structure
 - **Process-control kernel data structure**: PCB, kernel stacks
- Create PCBs of all processes
 - allocate stacks
 - privilege level setting using CONTROL register
 - Exception stack frame creation for new processes

Git Submission

- Tag your commit with “p2-submit”

References

1. Dasiewicz, Paul, A non-preemptive RTX Design Documentation
2. LPC17xx User's Manual
3. ARM Compilation Tools Version 5.0 Developer Guide
4. Software Interface Standard for Arm Cortex-based Microcontrollers, CMSIS Version 5.7.0

Acknowledgement

Slides modified from

- P. Dasiewicz
- J.C. Petkovich



Thank you!

Electrical and Computer Engineering Department